

uC/OS-II: Kernel Structure

Embedded OS Implementation

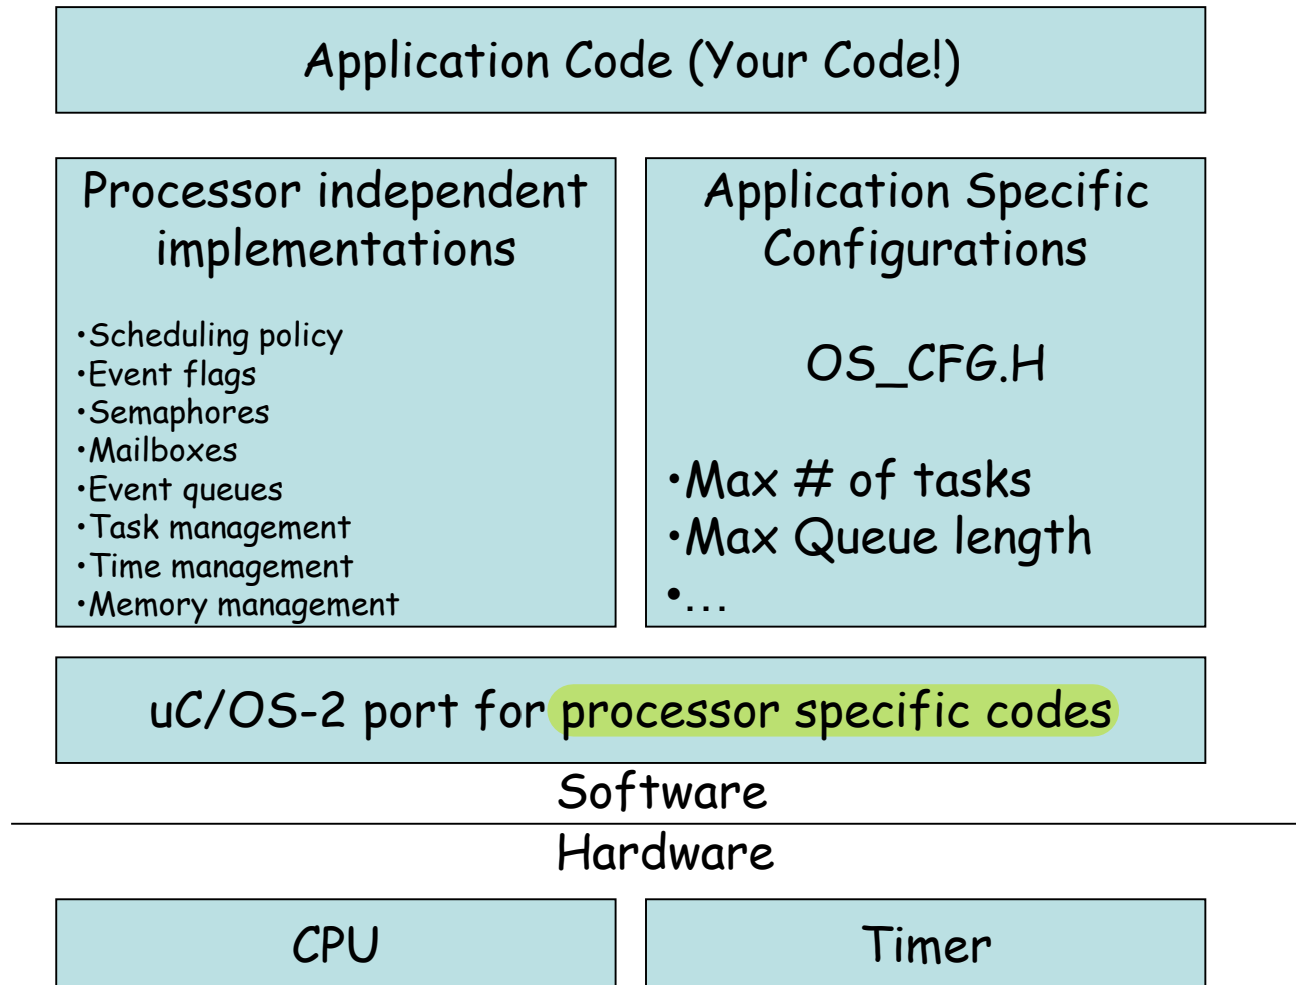
Prof. Ya-Shu Chen

National Taiwan University
of Science and Technology

Objectives

- To understand what a task is
- To learn how uC/OS-2 manages tasks

The uC/OS-2 File Structure



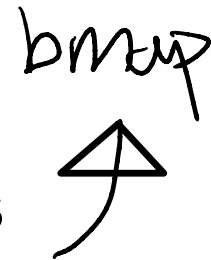
Tasks

- A task is an active entity which could do some computations
- In real-time systems, a periodic task is typically an infinite loop

```
void YourTask (void *pdata)           (1)
{
    for (;;) {                         (2)
        /* USER CODE */
        Call one of uC/OS-II's services:
        OSMboxPend();
        OSQPend();
        OSSemPend();
        OSTaskDel(OS_PRIO_SELF);
        OSTaskSuspend(OS_PRIO_SELF);
        OSTimeDly();
        OSTimeDlyHMSM();
        /* USER CODE */
    }
}
```

Delay itself for next event/period, so that other tasks can run.

Tasks



- uC/OS-2 can have up to 64 priorities
 - Each task must associate with an unique priority
 - 63 and 62 are reserved (idle, stat)
- Insufficient number of priority will damage the schedulability of a real-time scheduler
 - Tie-breaking for tasks having the same priority is an issue
 - Fortunately, many embedded systems have a limited number of tasks to run

Tasks

- A task is created by `OSTaskCreate()` or `OSTaskCreateExt()`
- The priority of a task can be changed by `OSTaskChangePrio()`
- A task could delete itself when done.

```
void YourTask (void *pdata)
{
    /* USER CODE */
    OSTaskDel(OS_PRIO_SELF);
}
```

The priority of the
current task

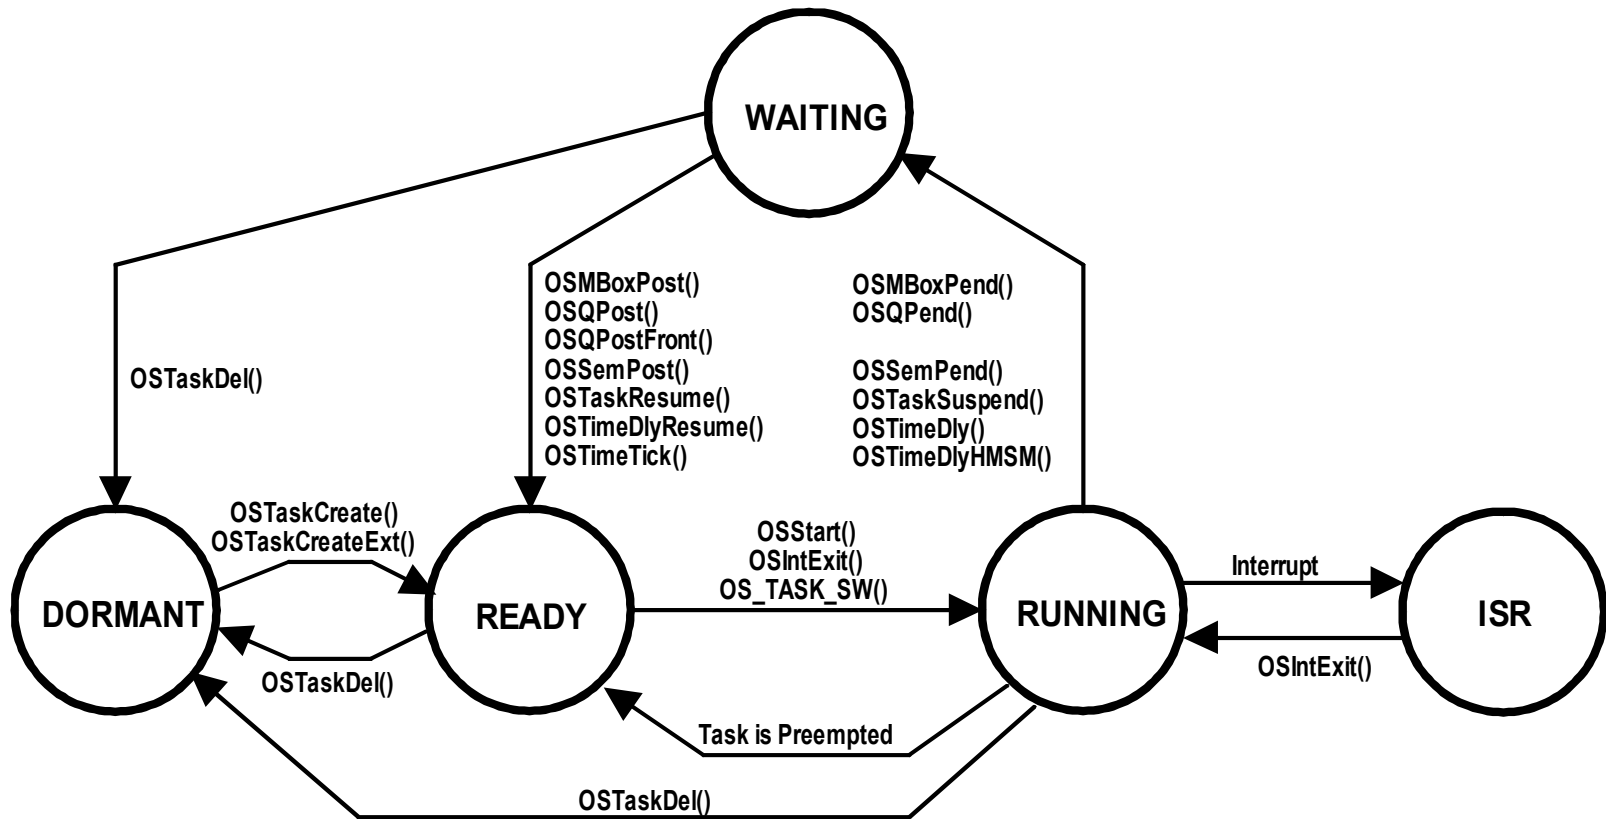
Task States

- Dormant: Procedures residing on RAM/ROM is not a task yet unless you call OSTaskCreate() to create one to execute them
- Ready: A task is neither delayed nor waiting for any event to occur
 - A task is ready once it is created
- Running: A ready task is running on the CPU
 - There must be only one running task.
 - The task running might be preempted and then become ready

Task States

- Waiting: A task is waiting for certain events to occur
 - Timer expiration, signaling of semaphores, messages in mailboxes, and etc
- ISR: A task is preempted by an interrupt
 - The stack of the interrupted task is utilized by the ISR

Task States



Task States

- A task can delay itself by calling `OSTimeDly()` or `OSTimeDlyHMSM()`.
 - The task is placed in the waiting state.
 - The task will be made ready by `OSTimeTick()`.
 - It is the clock ISR, you don't have to call it explicitly from your code.
- A task can wait for an event by `OSFlagPend()`, `OSSemPend()`, `OSMboxPend()`, or `OSQPend()`.
 - The task remains waiting until the occurrence of the desired event. (or timeout)
- The running task is always preempted by ISR's, unless interrupts are disabled.
 - ISR's could make one or more tasks ready by signaling events.
 - On the return of an ISR, the scheduler will check if rescheduling is needed.
- Once new tasks become ready, the next highest priority ready task is scheduled to run (due to occurrences of events, timer expirations).
- If no task is running and all tasks are not in the ready state, the idle task executes.

Task Control Blocks (TCB)

- A TCB is a RAM-resident per-task data structure
- Each task is associated with a TCB
 - All valid TCB's are doubly linked
 - Free TCB's are linked in a free list
- The contents of a TCB is saved/restored when a context-switch occurs
 - Task priority, delay counter, event to wait, location of the stack
 - *CPU registers are stored in the stack rather than in the TCB*

```

typedef struct os_tcb {
    OS_STK          *OSTCBStkPtr;
    #if OS_TASK_CREATE_EXT_EN
        void          *OSTCBExtPtr;
        OS_STK          *OSTCBStkBottom;
        INT32U          OSTCBStkSize;
        INT16U          OSTCBOpt;
        INT16U          OSTCBIId;
    #endif
    struct os_tcb *OSTCBNext;
    struct os_tcb *OSTCBPrev;
    #if (OS_Q_EN && (OS_MAX_QS >= 2)) || OS_MBOX_EN || OS_SEM_EN
        OS_EVENT          *OSTCBEvtPtr;
    #endif
    #if (OS_Q_EN && (OS_MAX_QS >= 2)) || OS_MBOX_EN
        void          *OSTCBMsg;
    #endif

    INT16U          OSTCBDly;
    INT8U           OSTCBStat;
    INT8U           OSTCBPrio;
    INT8U           OSTCBX;
    INT8U           OSTCBY;
    INT8U           OSTCBBitX;
    INT8U           OSTCBBitY;
    #if OS_TASK_DEL_EN
        BOOLEAN          OSTCBDelReq;
    #endif
} OS_TCB;

```

stack

⇒ double link

}

}

⇒ delete task

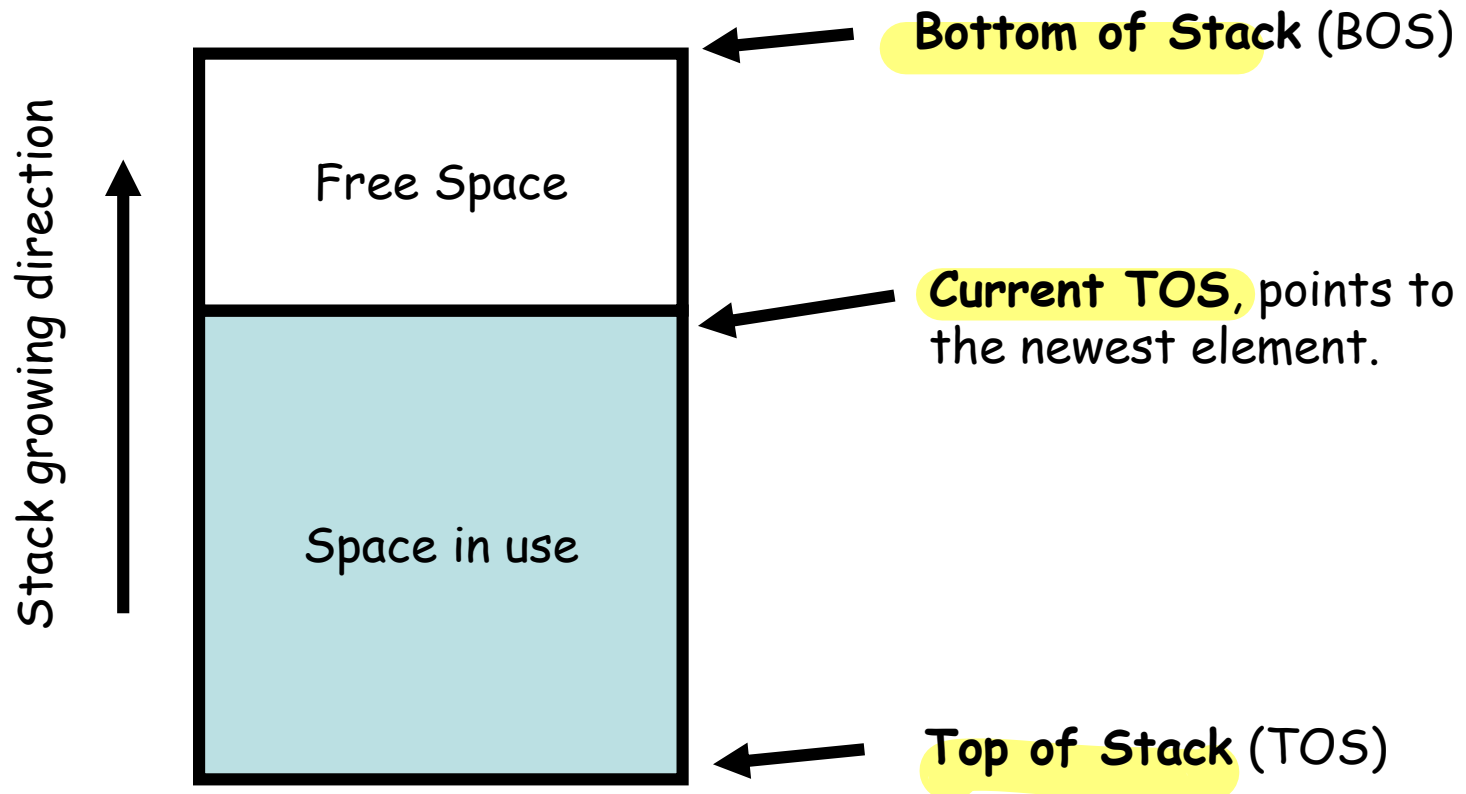
Task Control Blocks (TCB)

- **.OSTCBStkPtr** points to the **current** TOS for the task
 - It is the first entry of TCB so that it can be accessed directly via assembly language (offset=0)
- **.OSTCBExtPtr** is a pointer to a user-**definable** task control block extension.
 - Set OS_TASK_CREATE_EXT_EN to 1.
 - The pointer is set when OSTaskCreateExt() is called

Task Control Blocks (TCB)

- **.OSTCBStkBottom** is a pointer to the bottom of the task's stack
- **.OSTCBStkSize** holds the size of the stack in number of elements instead of bytes
 - The element size is the macro OS_STK.
 - Total stack size is OSTCBStkSize*OS_STK bytes
 - .OSTCBStkBottom and .OSTCBStkSize are used to check stack

Task Control Blocks (TCB)



Task Control Blocks (TCB)

- **.OSTCBOpt** holds “options” that can be passed to OSTaskCreateExt()
 - OS_TASK_OPT_STK_CHK: stack checking is enable for the task being created.
 - OS_TASK_OPT_STK_CLR: indicates that the stack needs to be cleared when the task is created.
 - OS_TASK_OPT_SAVE_FP: tells OSTaskCreateExt() that the task will be doing floating-point computations. Floating point processor’s registers must be saved to the stack on context-switches.
- **.OSTCBId**: holds an identifier for the task.
- **.OSTCBNext** and **.OSTCBPrev** are used to double link OS_TCBs
- **.OSTCBEVEventPtr** is pointer to an event control block.
- **.OSTCBMsg** is a pointer to a message that is sent to a task.
- **.OSTCBFlagNode** is a pointer to a flagnode.
- **.OSTCBFlagsRdy** maintains which event flags make the task ready.
- **.OSTCBDly** is used when:
 - a task needs to be delayed for a certain number of clock ticks, or
 - a task needs to pend for an event to occur with a timeout.
- **.OSTCBStat** contains the state of the task. (0 is ready to run)
- **.OSTCBPrio** contains the task priority.

Task Control Blocks (TCB)

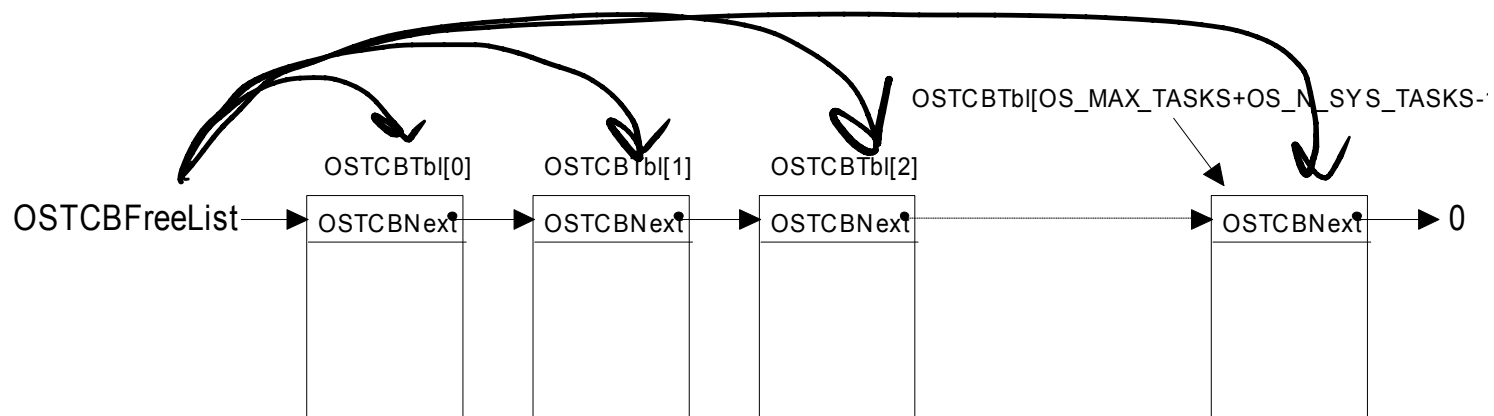
- **.OSTCBX .OSTCBY .OSTCBBitX and .OSTCBBitY**
 - They are used to accelerate the process of making a task ready to run or make a task wait for an event.

```
OSTCBY  = priority >> 3;  
OSTCBBitY      = OSMaPtbl[priority >> 3];  
OSTCBX  = priority & 0x07;  
OSTCBBitX      = OSMaPtbl[priority & 0x07];
```

- **.OSTCBDeIReq** is boolean used to indicate whether or not a task request that the current task to be deleted.
- **OS_MAX_TASKS** is specified in **OS_CFG.H**
 - # OS_TCBs allocated by µC/OS-II
- **OSTCBtbl[]** : where all OS_TCBs are placed.
- When µC/OS-II is initialized, all OS_TCBs in the table are linked in a singly linked list of free OS_TCBs

Task Control Blocks (TCB)

- A task receives/frees its OS_TCB from/to the free list
- An OS_TCB is initialized by the function OS_TCBInit(), which is called by OSTaskCreate().



```

INT8U OS_TCBInit (INT8U prio, OS_STK *ptos, OS_STK *pbos, INT16U id, INT32U stk_size, void *pext, INT16U
opt)
{
    #if OS_CRITICAL_METHOD == 3
        OS_CPU_SR cpu_sr;
    #endif
    OS_TCB *ptcb;

    OS_ENTER_CRITICAL();
    ptcb = OSTCBFreeList;
    if (ptcb != (OS_TCB *)0) {
        OSTCBFreeList = ptcb->OSTCBNext;
    }
    OS_EXIT_CRITICAL();

    ptcb->OSTCBStkPtr = ptos;
    ptcb->OSTCBPrio = (INT8U)prio;
    ptcb->OSTCBStat = OS_STAT_RDY;
    ptcb->OSTCBDly = 0;

    #if OS_TASK_CREATE_EXT_EN > 0
        ptcb->OSTCBExtPtr = pext;
        ptcb->OSTCBStkSize = stk_size;
        ptcb->OSTCBStkBottom = pbos;
        ptcb->OSTCBOpt = opt;
        ptcb->OSTCBIId = id;
    #else
        pext = pext;
        stk_size = stk_size;
        pbos = pbos;
        opt = opt;
        id = id;
    #endif

    #if OS_TASK_DEL_EN > 0
        ptcb->OSTCBDelReq = OS_NO_ERR;
    #endif

    ptcb->OSTCBY = prio >> 3;
    ptcb->OSTCBBitY = OSMaPtbl[ptcb->OSTCBY];
    ptcb->OSTCBX = prio & 0x07;
    ptcb->OSTCBBitX = OSMaPtbl[ptcb->OSTCBX];

```

Get a free TCB from
the free list

清空

```

#if OS_EVENT_EN > 0
    ptcb->OSTCBEventPtr = (OS_EVENT *)0;          /* Task is not pending on an event */
#endif

#if (OS_VERSION >= 251) && (OS_FLAG_EN > 0) && (OS_MAX_FLAGS > 0) && (OS_TASK_DEL_EN > 0)
    ptcb->OSTCBFlagNode = (OS_FLAG_NODE *)0;      /* Task is not pending on an event flag */
#endif

#if (OS_MBOX_EN > 0) || ((OS_Q_EN > 0) && (OS_MAX_QS > 0))
    ptcb->OSTCBMsg      = (void *)0;              /* No message received */
#endif

#if OS_VERSION >= 204
    OSTCBInitHook(ptcb);
#endif

    OSTaskCreateHook(ptcb);                      /* Call user defined hook */

    OS_ENTER_CRITICAL();
    OSTCBPrioTbl[prio] = ptcb;
    ptcb->OSTCBNext     = OSTCBList;
    ptcb->OSTCBPrev     = (OS_TCB *)0;
    if (OSTCBList != (OS_TCB *)0) {
        OSTCBList->OSTCBPrev = ptcb;
    }
    OSTCBList           = ptcb;
    OSRdyGrp            |= ptcb->OSTCBBitY;
    OSRdyTbl[ptcb->OSTCBY] |= ptcb->OSTCBBitX;
    OS_EXIT_CRITICAL();
    return (OS_NO_ERR);
}
OS_EXIT_CRITICAL();
return (OS_NO_MORE_TCB);
}

```

User-defined hook is called here.

Priority table

TCB list

/* Link into TCB chain

Ready list

/* Make task ready to run

Critical Sections

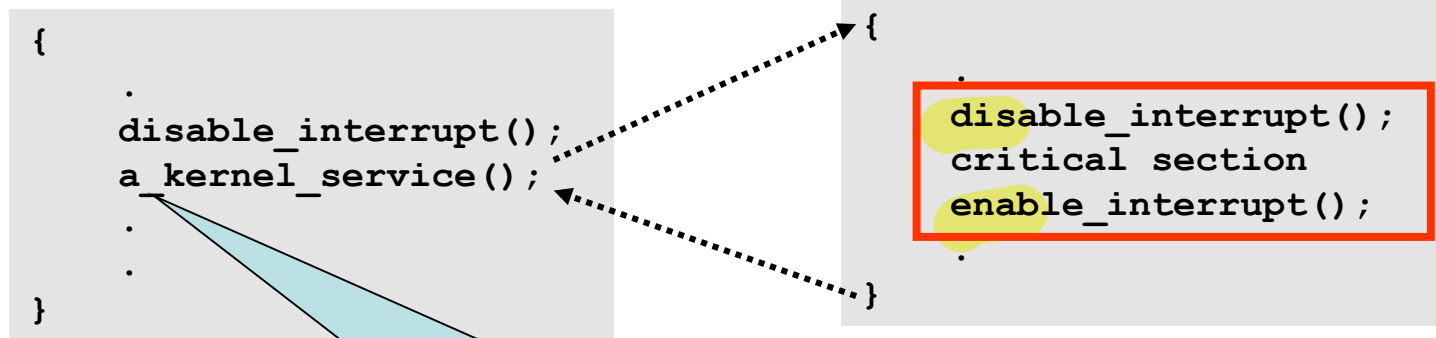
- A critical section is a portion of code that is not safe from race conditions
- They can be protected by interrupt disabling/enabling interrupts or semaphores
 - However, the use of semaphores imposes much more overheads than enabling/disabling interrupts
 - A RTOS kernel itself mostly use interrupts disabling/enabling to protect critical sections
- Once interrupts are disabled, neither context switches nor any other ISR's can occur

Critical Sections

- The states of the processor must be carefully maintained across multiple calls of `OS_ENTER_CRITICAL()` / `OS_EXIT_CRITICAL()`
- There are three possible implementations for the maintenance of process states:
 - Interrupt `enabling/disabling` instructions
 - Interrupt `status save/restore` onto/from stacks
 - `Processor Status Word (PSW)` save/restore onto/from memory variables

Critical Sections

- OS_CRITICAL_METHOD=1
- Interrupt enabling/disabling instructions.
- The simplest way, however, this approach does not have the sense of “save” and “restore”



Interrupts are now
implicitly re-enabled!

ISR → ~~PSW~~ → CPU state X ⇒ PSW

Critical Sections

- OS_CRITICAL_METHOD=2
- Processor Status Word (PSW) can be saved/restored onto/from stacks
 - PSW's of nested interrupt enable/disable operations can be exactly recorded in stacks

```
#define OS_ENTER_CRITICAL() \  
    asm("PUSH    PSW");  
    asm("DI");  
  
#define OS_EXIT_CRITICAL() \  
    asm("POP     PSW");
```


Task Scheduling

- The scheduler always schedules the highest-priority ready task to run
- Task-level scheduling and ISR-level scheduling are done by `OS_Sched()` and `OSIntExit()`, respectively
- uC/OS-2 scheduling time is a predictable amount of time, i.e., a constant time
 - For example, the design of the ready list intends to achieve this objective



```
void OSSched (void)
{
    INT8U y;
    OS_ENTER_CRITICAL();
    if ((OSLockNesting | OSIntNesting) == 0) {           (1)
        y = OSUnMapTbl[OSRdyGrp];                       (2)
        OSPrioHighRdy = (INT8U)((y << 3) + OSUnMapTbl[OSRdyTbl[y]]); (2)
        if (OSPrioHighRdy != OSPrioCur) {              (3)
            OSTCBHighRdy = OSTCBPrioTbl[OSPrioHighRdy]; (4)
            OSCtxSwCtr++;                                (5)
            OS_TASK_SW();                                (6)
        }
    }
    OS_EXIT_CRITICAL();
}
```

- (1) Rescheduling will not be performed if the scheduler is locked or some interrupt is currently serviced (why?).
- (2) Find the highest-priority ready task.
- (3) If it is not the current task, then
- (4) ~ (6) Perform a context-switch.

OS_TASK_SW() is a macro: "asm int 0x80"



```
void OSIntExit (void)
{
    OS_ENTER_CRITICAL();
    if ((--OSIntNesting | OSLockNesting) == 0) {
        OSIntExitY = OSUnMapTbl[OSRdyGrp];
        OSPrioHighRdy = (INT8U)((OSIntExitY << 3) +
                                OSUnMapTbl[OSRdyTbl[OSIntExitY]]);
        if (OSPrioHighRdy != OSPrioCur) {
            OSTCBHighRdy = OSTCBPrioTbl[OSPrioHighRdy];
            OSCtxSwCtr++;
            OSIntCtxSw();
        }
    }
    OS_EXIT_CRITICAL();
}
```

If scheduler is not locked and no interrupt nesting

If there is another high-priority task ready

A context switch is performed.

Note that **OSIntCtxSw()** is called instead of calling **OS_TASK_SW()** because the **ISR** already saves the **CPU registers** onto the stack.

```
void OSIntEnter (void)
{
    OS_ENTER_CRITICAL();
    OSIntNesting++;
    OS_EXIT_CRITICAL();
}
```

只有一半

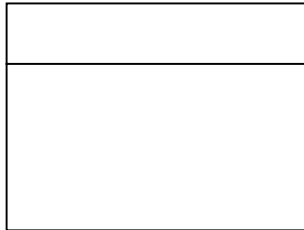
Task Level Context Switch

- By default, context switches are handled at interrupt-level, therefore task-level scheduling will invoke a software interrupt to simulate that
 - Hardware dependent, porting must be done

Low Priority Task

OS_TCB

OSTCBCur →



Low Memory



High Memory

CPU

• SP

R4

R3

R2

R1

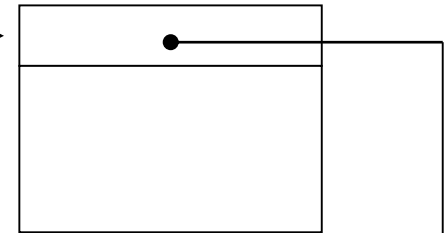
PC

PSW

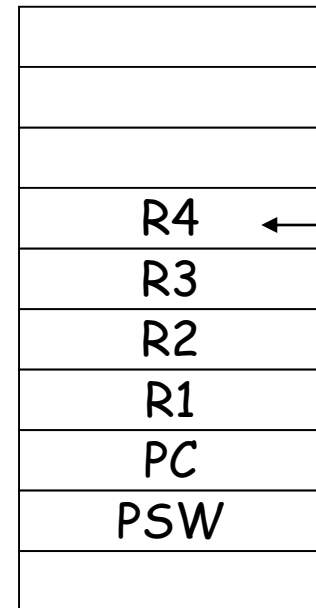
High Priority Task

OS_TCB

OSTCBHighRdy →

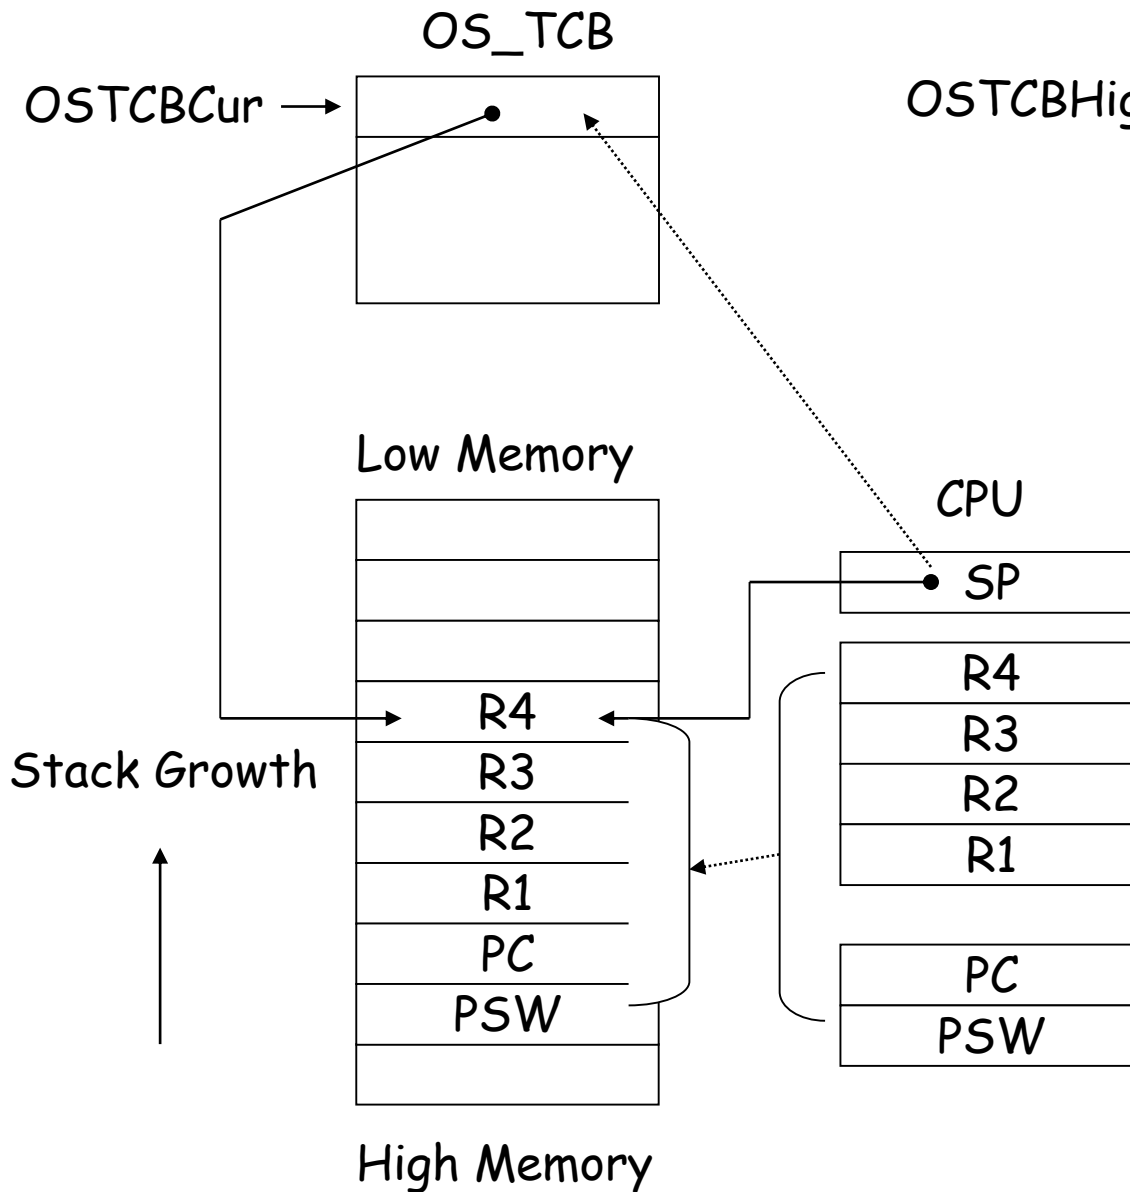


Low Memory

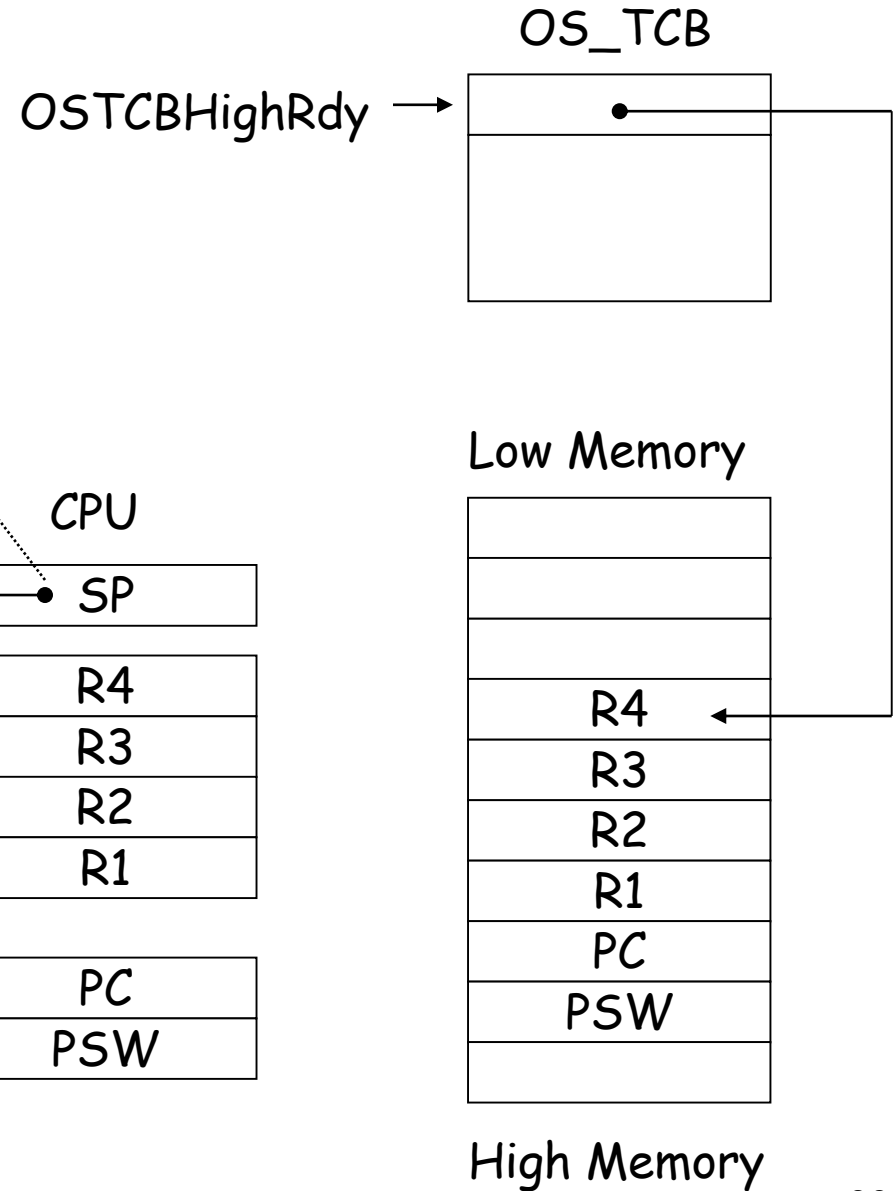


High Memory

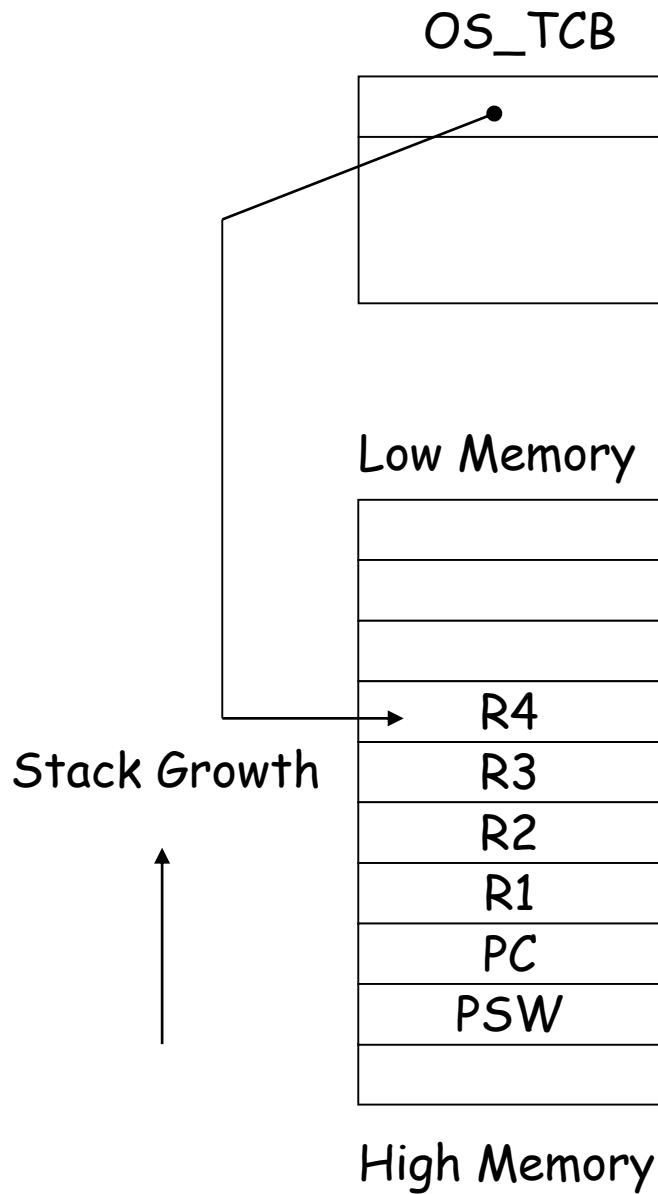
Low Priority Task



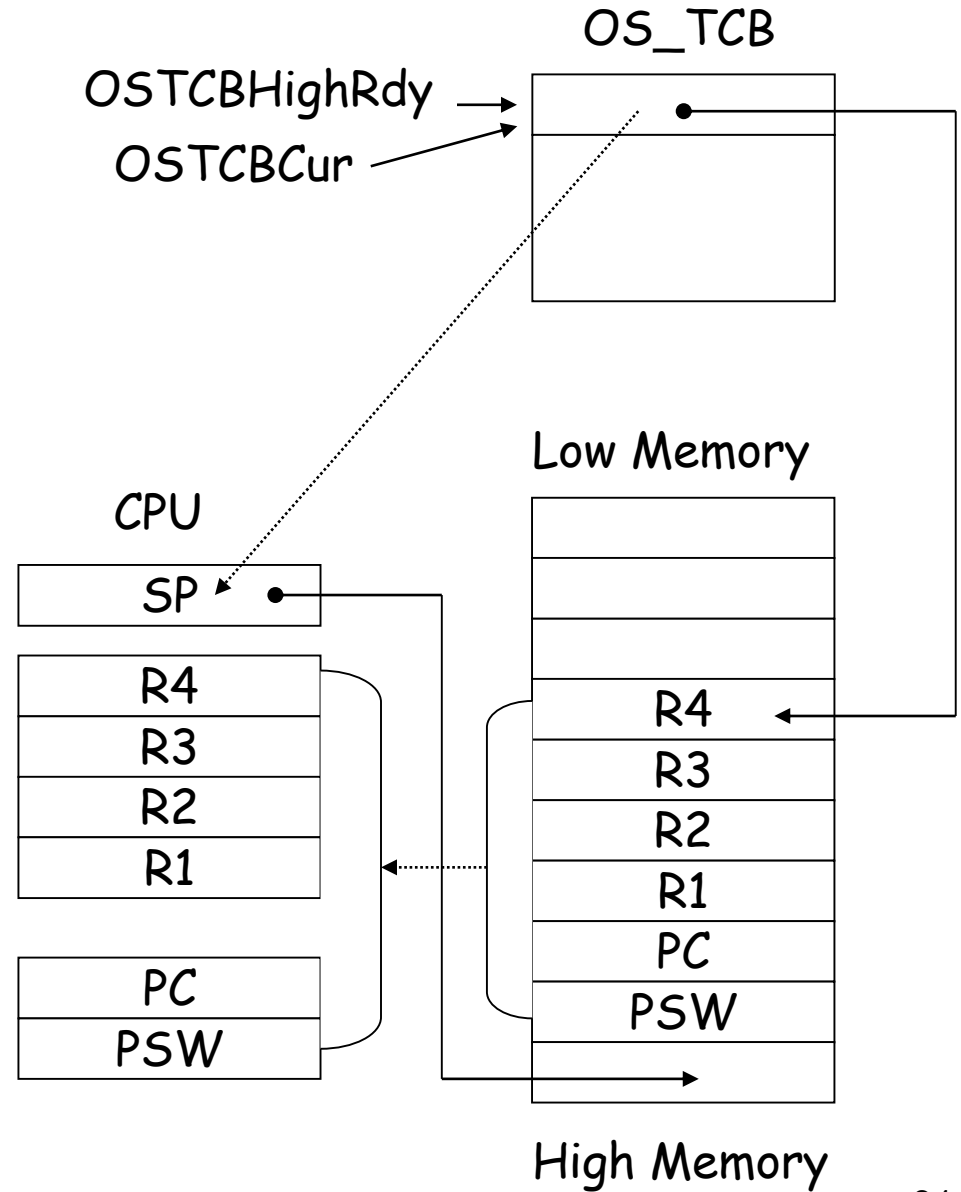
High Priority Task



Low Priority Task



High Priority Task



Locking and Unlocking the Scheduler

- `OSSchedLock()` prevent high-priority ready tasks from being scheduled to run while interrupts are still recognized
- `OSSchedLock()` and `OSSchedUnlock()` are used in pairs
- `OSLockNesting` keeps track of the number of `OSSchedLock()` has been called

Locking and Unlocking the Scheduler

- After calling `OSSchedLock()`, you must not call kernel services which might cause context switch, such as `OSFlagPend()`, `OSMboxPend()`, `OSMutexPend()`, `OSQPend()`, `OSSemPend()`, `OSTaskSuspend()`, `OSTimeDly`, `OSTimeDlyHMSM()` until `OSLockNesting == 0`. Or the system will be locked up
- To lock the scheduler is to prevent from race conditions while interrupts can still be handled

OSSchedLock()

```
void OSSchedLock (void)
{
    #if OS_CRITICAL_METHOD == 3          /* Allocate storage for CPU status register */
        OS_CPU_SR  cpu_sr;
    #endif

    if (OSRunning == TRUE) {             /* Make sure multitasking is running */
        OS_ENTER_CRITICAL();
        if (OSLockNesting < 255) { /* Prevent OSLockNesting from wrapping back to 0 */
            OSLockNesting++;         /* Increment lock nesting level */
        }
        OS_EXIT_CRITICAL();
    }
}
```

OSSchedUnlock()

```
void OSSchedUnlock (void)
{
    #if OS_CRITICAL_METHOD == 3                /* Allocate storage for CPU status register */
        OS_CPU_SR  cpu_sr;
    #endif

    if (OSRunning == TRUE) {                  /* Make sure multitasking is running */
        OS_ENTER_CRITICAL();
        if (OSLockNesting > 0) {              /* Do not decrement if already 0 */
            OSLockNesting--;                  /* Decrement lock nesting level */
            if ((OSLockNesting == 0) &&
                (OSIntNesting == 0)) {        /* See if sched. enabled and not an ISR */
                OS_EXIT_CRITICAL();
                OS_Sched();                   /* See if a HPT is ready */
            } else {
                OS_EXIT_CRITICAL();
            }
        } else {
            OS_EXIT_CRITICAL();
        }
    }
}
```

The Idle Task

- The idle task is always the lowest-priority task and can not be deleted or suspended
- To conserve power dissipation, you can issue a HALT instruction in the idle task

```
void OS_TaskIdle (void *pdata)
{
    #if OS_CRITICAL_METHOD == 3
        OS_CPU_SR cpu_sr;
    #endif

    pdata = pdata;
    for (;;) {
        OS_ENTER_CRITICAL();
        OSIdleCtr++;
        OS_EXIT_CRITICAL();
        OSTaskIdleHook();
    }
}
```

Summary

- In this class, you should learn:
 - What a task is, how uC/OS-2 manages a task, and related data structures
 - How the scheduler works in uC/OS-2
 - The responsibility of the idle task