

# Independent Task Scheduling

Embedded OS Implementation

Prof. Ya-Shu Chen

National Taiwan University  
of Science and Technology


# Motivation

- Take yourself as an example
  - Naturally you have a number of things to do with time pressure
    - Project deadlines, meeting time, class time, and deadlines for bills
  - Some of them regularly recur but some don't
    - To eat meal on 12:30 everyday
    - Go to the movies on 8:00pm


避免 overload

# Motivation

- You schedule yourself to meet deadlines
  - Course A: one homework is announced every 9 days, each costs you 4.5 days to do
  - Course B: one homework is announced every 6 days, each costs you 3 days to do
- You miss deadlines of one course if your policy favors either one course




Favor course B



Favor course A


# Motivation

- **Schedule to meet deadlines** (cont'd.) *change*
  - Course A: (4.5, 9) *0.5/day*
  - Course B: (3, 6) *0.5/day*
- All deadlines are met if you whatever has the **closest deadline** *A & B*



# Motivation

- You schedule yourself to survive overloads
  - $(1,2), (1,3), (1,1.5) \Rightarrow (3, 2, 4, 6) \rightarrow \text{overload}$   
 $0.5 \quad 0.37 \quad 0.63 \Rightarrow 1.5 > 1 \text{ day} \Rightarrow \text{impossible}$



Favoring  $(1,1.5) \rightarrow (1,3) \rightarrow (1,2)$   
 2 lecturers are happy, 1 will flunk you though...


Do whatever has the closest deadline.  
*You are in deep shit!*

# Cyclic-Executive

# Cyclic Executive


一週化 ⇒ 照表操課

- The system repeatedly exercises a static schedule
  - A table-driven approach
- Many existing systems still take this approach
  - Easy to debug and easy to visualize
    - Highly deterministic
  - Hard to program, to modify, and to upgrade
    - A program should be divided into many pieces (like an FSM)



# Cyclic Executive

- The table emulates an infinite loop of routines
  - However, a single independent loop is not enough to many complicated systems
  - Multiple concurrent loops should be considered
- How large should the table be when there are multiple loops?





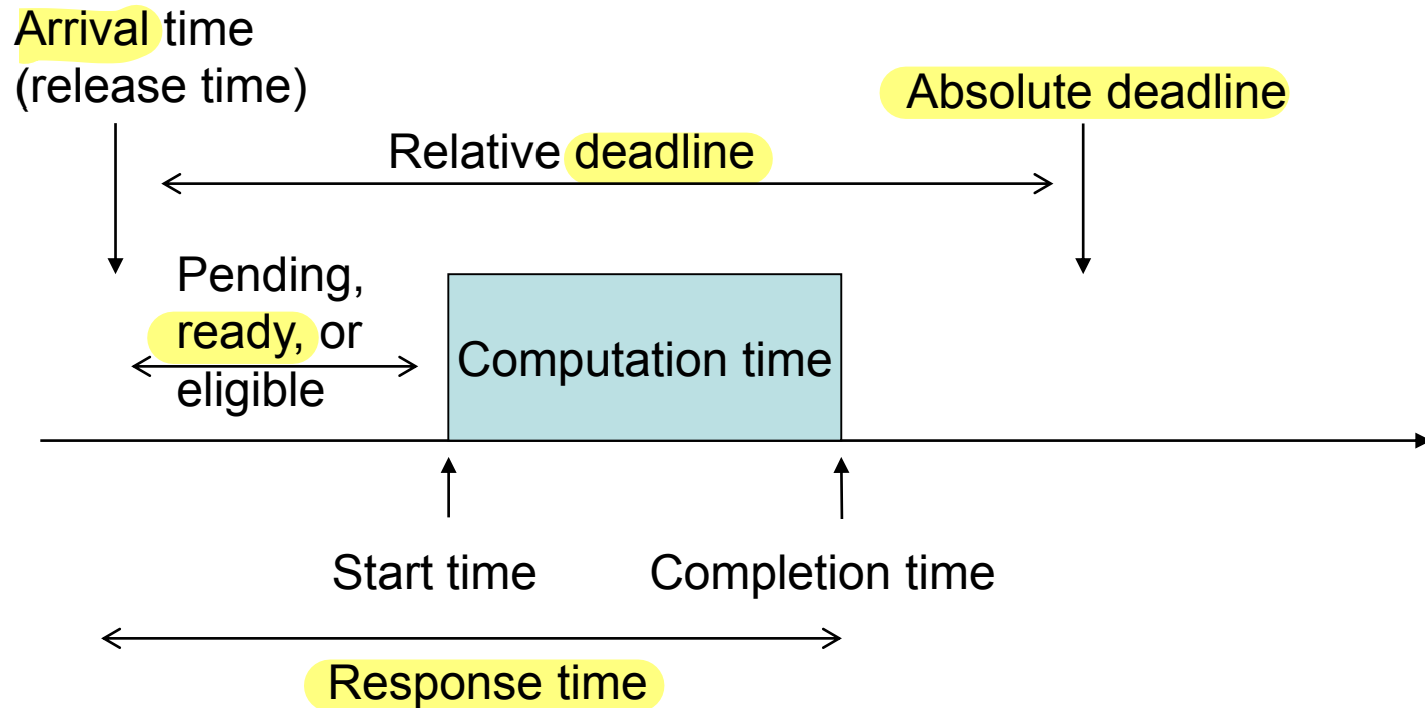
# Cyclic Executive

- **Definition:** Let the **hyper-period** of a collection of loops be a time interval which's length is the **least-common-multiplier of the loops' lengths**
  - Let the length of the hyper-period be abbreviated as **"h"** 最小公倍數 (工作週期)
- **Theorem:** The number of routines **to be executed** in any time interval **[t,t+x]** is identical to that in **[t+h,t+h+x]**

# Priority-Driven Scheduling

# System Model

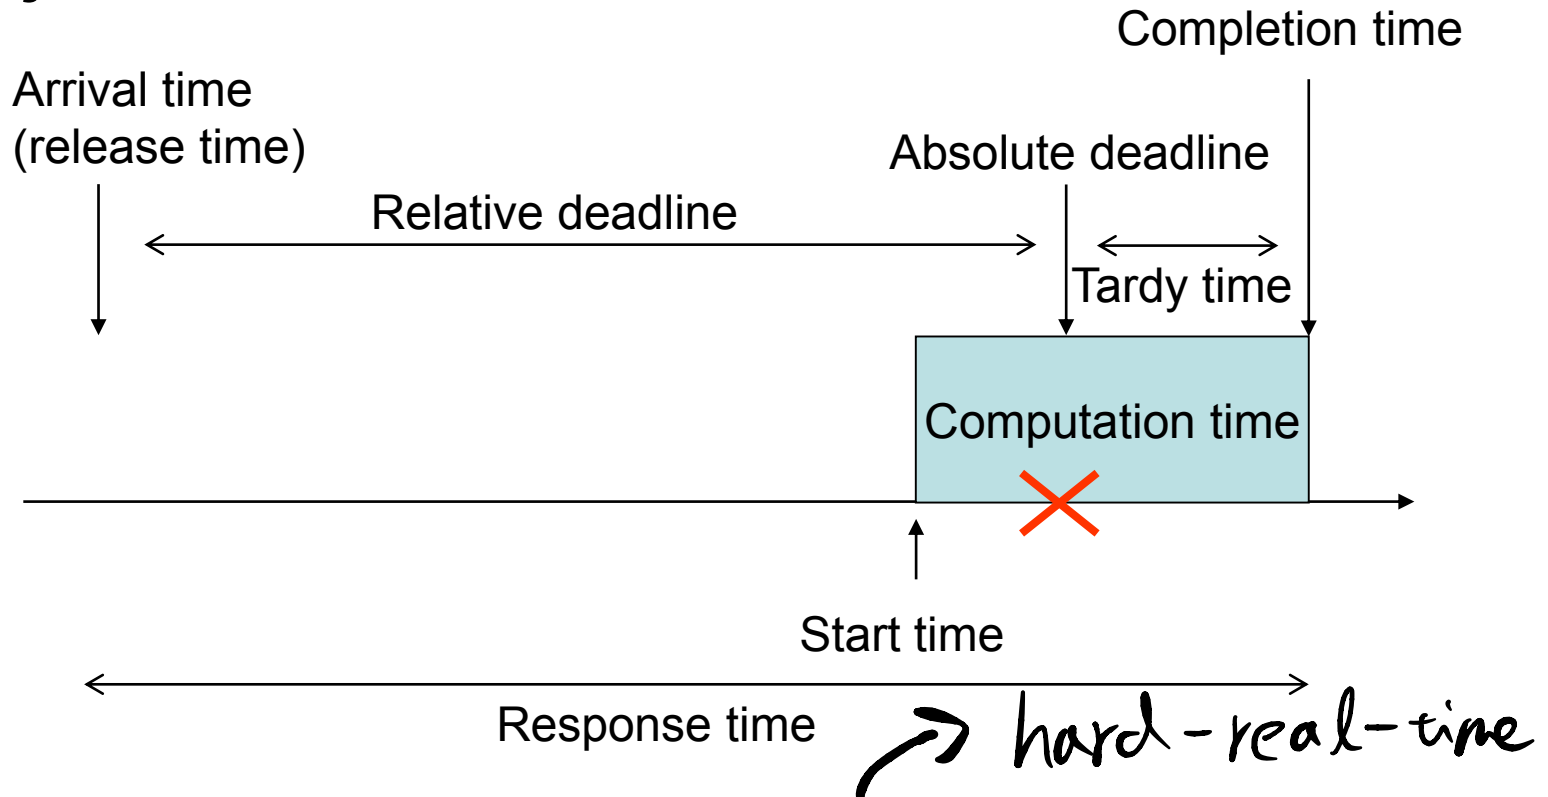
- A job with real-time constraints



The job completes before its deadline, that means the deadline is satisfied.

# System Model

- A job with real-time constraints




The job completes **after its deadline**, that means the deadline is **violated** or an **overflow** occurs.

# System Model

- A task set is of a number of tasks
  - $\{T_1, T_2, \dots, T_n\}$
  - Tasks share nothing but CPU, and tasks are independent to one another
- A task  $T_i$  is a template of jobs, where jobs refer to recurrences of task.
  - Every job executes the same piece of code
    - Of course, different input and run-time conditions cause jobs behaves differently
    - $J_{i,j}$  refers to the j-th job of task  $T_i$
  - The computation time  $C_i$  of jobs is bounded and known a priori

# System Model

- A purely periodic task
  - Jobs of a task **T** recur every fixed time interval **p**
  - A job must be completed before the next job arrives
    - Relative deadlines for jobs are, implicitly, the period
  - T is defined as (c,p) 固定的週期的



Periodic task  $T_1=(3,6)$

# System Model

- **Priority**
  - Reflect the urgency of jobs
  - Any job inherits its task's priority
- **Preemptivity** *first in, first out*
  - As a high-priority task arrives, it preempts the execution of any low-priority tasks

# System Model


- Checklist
  - Periodic tasks
  - Real-time constraints
  - Priority
  - Preemptivity



# Rate-Monotonic Scheduling

- Task-level fixed-priority scheduling
  - All jobs inherit its task's priority
  - Usually abbreviated as fixed-priority scheduling
- Tasks' priorities are inversely proportional to their period lengths

反比




# Rate-Monotonic Scheduling

✓ 同時到最長的 Response Time

- **Critical instant** (critical instance) of task  $T_i$ 
  - A job  $J_{i,c}$  of task  $T_i$  released at  $T_i$ 's critical instant would have the **longest response time**
  - $J_{i,c}$  would be the one that is **"hardest"** to meet its deadline
  - If  $J_{i,c}$  succeeds in satisfying its deadline, then any job of  $T_i$  always succeeds for any cases
    - Since in any other cases deadlines are easier to meet

# Rate-Monotonic Scheduling

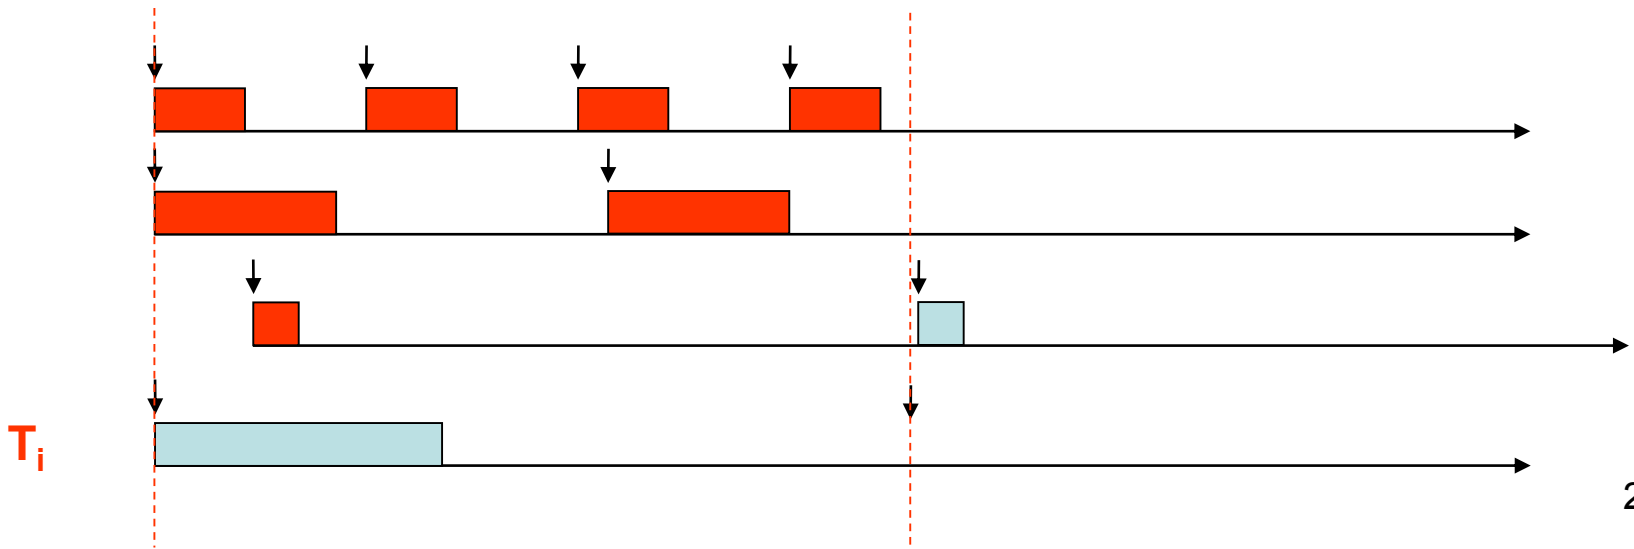
- **Theorem:** A critical instant of any task  $T_i$  occurs when one of its job  $J_{i,c}$  is released at the same time with a job of every higher-priority task (i.e., in-phase).



# Rate-Monotonic Scheduling


- Proof: “interferences” from high-priority tasks is the maximum within the first period of  $T_i$

$$\sum_{j < i} c_j \left\lceil \frac{p_i}{p_j} \right\rceil$$



# Rate-Monotonic Scheduling

- Critical instant: Why ceiling function?




# Rate-Monotonic Scheduling

- Response time analysis
  - The response time of the job of  $T_i$  at critical instant can be calculated by the following recursive function

$$r_0 = \sum_{\forall i} c_i$$

$$r_n = \sum_{\forall i} c_i \left\lceil \frac{r_{n-1}}{p_i} \right\rceil$$



- Observation: the sequence of  $r_x$ ,  $x \geq 0$  may or may not converge

$$r_n = r_{n-1} \Rightarrow \text{收敛}$$

# Rate-Monotonic Scheduling

- **Theorem:** Given a task set  $\{T_1, T_2, \dots, T_n\}$ , if at critical instant the response time of the first job of task  $T_i$ , for each  $i$ , converges no later than  $p_i$ , then jobs never miss their deadlines
- Observations
  - If the task set survives critical instant, then it will survive any task phasing
  - The analysis is an exact schedulability test for RMS
  - Usually referred to as “Rate-Monotonic Analysis”, RMA for short

# Rate-Monotonic Scheduling

- Example:  $T1=(2,5)$ ,  $T2=(2,7)$ ,  $T3=(3,8)$

– T1:

- $R_0=2 \leq 5$  ok

– T2:

- $R_0=2+2=4 \leq 7$

- $R_1=2 * \lceil 4/5 \rceil + 2 * \lceil 4/7 \rceil = 4 \leq 7$  ok

$R_1 = R_0$   
 $\Rightarrow$  收敛

– T3:  $2+2=4$

- $R_0=2+2+3=7 \leq 8$

- $R_1=2 * \lceil 7/5 \rceil + 2 * \lceil 7/7 \rceil + 3 * \lceil 7/8 \rceil = 9 > 8$  failed

– Note: each task succeeds  $\rightarrow$  the task set

$$4+2+3=9$$



# Rate-Monotonic Scheduling

- Proof:
  - If the response time converges at  $r_n$ , then the first lowest-priority job completes at  $r_n$
  - If  $r_n$  is before  $p_n$ , then the first lowest-priority job meets its deadline if critical instant occurs
  - Since the job survives critical instant, it always succeed satisfying its deadline under any task phasing

# Rate-Monotonic Scheduling

- Test every  $T_i$  for schedulability!!
  - $\{T1=(3,6), T2=(3.1,9), T3=(1,18)\}$
  - Response analysis of  $T3$ :
    - $R0=7.1, R1=10.1, R2=13.2, R3=16.2, R4=16.2 < 18$
  - $\{T1, T2, T3\}$  is schedulable!?
  - However,  $\{T1, T2\}$  is not schedulable!!!

$$T_1: R_0 = 3 \leq 6$$

$$T_2: R_0 = 3 + 3.1 = 6.1 \leq 9$$

$$R_1 = 3 \times \left\lceil \frac{6.1}{6} \right\rceil + 3.1 \times \left\lceil \frac{6.1}{9} \right\rceil$$

$$= 6 + 3.1 = 9.1 > 9 \Rightarrow \text{not schedulable}$$

$T_3:$

$$R_0 = 3 + 3.1 + 1 = 7.1 \leq 18$$

$$R_1 = 3 \times \left\lceil \frac{7.1}{6} \right\rceil + 3.1 \times \left\lceil \frac{7.1}{9} \right\rceil + 1 \times \left\lceil \frac{7.1}{18} \right\rceil$$

$$= 6 + 3.1 + 1 = 10.1 \leq 18$$

$$R_2 = 3 \times \left\lceil \frac{10.1}{6} \right\rceil + 3.1 \times \left\lceil \frac{10.1}{9} \right\rceil + 1 \times \left\lceil \frac{10.1}{18} \right\rceil = 13.2 \leq 18$$

$$R_3 = 3 \times \left\lceil \frac{13.2}{6} \right\rceil + 3.1 \times \left\lceil \frac{13.2}{9} \right\rceil + 1 \times \left\lceil \frac{13.2}{18} \right\rceil$$

$$= 16.2 \leq 18$$

$$R_4 = 3 \times \left\lceil \frac{16.2}{6} \right\rceil + 3.1 \times \left\lceil \frac{16.2}{9} \right\rceil + 1 \times \left\lceil \frac{16.2}{18} \right\rceil$$

$$= 9 + 6.2 + 1 = 16.2 \leq 18$$

$$R_4 = R_3 \Rightarrow \text{4x 鈞}$$

# Rate-Monotonic Scheduling

- Computational complexity
  - $O(n^2 * p_n)$ , pseudo-polynomial time
  - Would be *extremely slow* when periods of tasks are small and prime to one another
  - Would be very fast when periods are harmonically related

# Rate-Monotonic Scheduling

- Phenomena
  - Even though RMA is an exact test for fixed-priority scheduling, it is not often used, especially for those dynamic systems
  - RMA is more suitable for static systems
  - Are there any schedulability tests being efficient enough for on-line implementation?
    - No slower than polynomial time

# Rate-Monotonic Scheduling

- A trivial schedulability test
  - The system accepts a task set  $T$  if the following conditions are both true
    - There are no other tasks in the system
    - $c_1/p_1 \leq 1$
  - The algorithm is efficient enough (i.e.,  $O(1)$ )
  - Too conservative!! Very Poor CPU utilization!!

# Rate-Monotonic Scheduling

- Definition

- Utilization factor of task  $T=(c,p)$  is defined as

$$\frac{c}{p}$$

- CPU utilization of a task set  $\{T_1, T_2, \dots, T_n\}$  is

$$U = \sum_{i=1}^n \frac{c_i}{p_i}$$

- Observation: if the total utilization exceeds 1 then the task set is not schedulable

# Rate-Monotonic Scheduling

- **Theorem:** [LL73] Given a task set  $\{T_1, T_2, \dots, T_n\}$ . It is schedulable by RMS if

$$\sum_{i=1}^n \frac{c_i}{p_i} \leq U(n) = n(2^{1/n} - 1)$$

- Observation:
  - If the test succeeds then the task is schedulable
  - A sufficient condition for schedulability



# Rate-Monotonic Scheduling

- Proof: Let us consider two tasks only



$$C_1 \leq P_2 - P_1(\lfloor P_2/P_1 \rfloor)$$

The largest possible C2 is

$$P_2 - C_1(\lceil P_2/P_1 \rceil)$$

Total utilization factor is


$$U = 1 + \underline{C_1(1/P_1 - (1/P_2)(\lceil P_2/P_1 \rceil))}$$

T2's 2nd job does not overlap the immediately preceding job of T1

- U monotonically decreases with  $C_1$ 
  - Because U never greater than 1, so the rightmost term in the last equation is always negative

# Rate-Monotonic Scheduling

- Proof: cont'd.



$$C_1 \geq P_2 - P_1(\lfloor P_2/P_1 \rfloor)$$

The largest possible  $C_2$  is

$$\underline{C_1}(\lfloor P_2/P_1 \rfloor) + P_1(\lfloor P_2/P_1 \rfloor)$$

Total utilization factor is

$$U = (P_1/P_2)\lfloor P_2/P_1 \rfloor + C_1((1/P_1) - (1/P_2)(\lfloor P_2/P_1 \rfloor))$$

T2's 2nd job overlaps the immediately preceding job of T1


- $U$  monotonically increases with  $C_1$

# Rate-Monotonic Scheduling

- Proof: Cont'd.

- It can be found that the minimal  $U$  occurs at

$$C_1 = P_2 - P_1(\lfloor P_2/P_1 \rfloor)$$




- By some differentiation, the minimal achievable utilization is

$$U(2) = 2(2^{1/2} - 1)$$

# Rate-Monotonic Scheduling

- To find “the smallest” among “the largest achievable processor utilizations that can be achieved by different task sets”




# Rate-Monotonic Scheduling

- Simon says: To generalize the proof to  $n$  tasks is easy :)
- If a task set of  $n$  tasks has a total utilization being no greater than  $U(n)$ , then it is guaranteed to be schedulable by RMS
  - Because the most hard-to-schedule task set having the same total utilization is schedulable
  - The test's time complexity is  $O(n)$ , which is very efficient for on-line implementation

# Rate-Monotonic Scheduling

- When  $x \rightarrow$  infinitely large,  $U(x) \rightarrow 0.68$




1	1
2	0.828427
3	0.779763
4	0.756828
5	0.743492
6	0.734772
7	0.728627
8	0.724062
9	0.720538
10	0.717735
11	0.715452
12	0.713557
13	0.711959
14	0.710593
15	0.709412

# Rate-Monotonic Scheduling


- Example 1: (1,3), (2,5)

– Utilization =  $0.73 \leq U(2) = 0.828$       $\frac{1}{3} + \frac{2}{5} = 0.73$




- Example 2: (4.5,9), (3,6)

– Utilization =  $100\% > U(2) = 0.828$       $\frac{1}{2} + \frac{1}{2} = 1$



# Rate-Monotonic Scheduling

- Example 3: (1,2), (2,4)  
– Utilization = 100% >  $U(2) = 0.828$   $\frac{1}{2} + \frac{2}{4} = 1$




- Example 2 and 3 shows that, we know nothing about those task sets of total utilization > the utilization bound!  $\hookrightarrow$  不一定
- But we do know those  $\leq$  the utilization bound is schedulable!



# Rate-Monotonic Scheduling

- Sufficiency but no necessity
  - Utilization test provides a fast way to check if a task set is schedulable
  - Any task set fails utilization test, does not implies that it is not schedulable



$U$ : universe of task sets

$\sim S$ : task sets unschedulable by RMS

- Example 2

$S$ : task sets schedulable by RMS

- Example 1 and Example 3

$A$ : Those can be found by utilization test

- Example 1

Example 3 is in  $S-A$

# Rate-Monotonic Scheduling

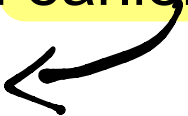
- Summary
  - Explicit prioritization over tasks
  - To **decide** task sets' schedulability is costly
  - Sufficient tests were developed for fast admission control

# Priority-Driven Scheduling: Dynamic-Priority Scheduling

# Earliest-Deadline-First Scheduling

- Definition
  - Feasible
    - A set of tasks is said to be feasible if there is some way to schedule the tasks without any deadline violations
  - Schedulable
    - Given a scheduling algorithm A
    - A set of tasks is said to be schedulable if algorithm A successfully schedule the tasks without any deadline violations
- Observations
  - A feasible task set may not be schedulable by RMS
  - If a task set is schedulable by some algorithm A, then it is feasible

# Earliest-Deadline-First Scheduling

- EDF scheduling algorithm always pick a pending job which has the earliest deadline for execution
  - A job having an earlier deadline is assigned to a higher “priority”
  - Priority in EDF is not a task-wide notion
    - Jobs of a task may have different “priorities”
    - But due to the relative deadline of a job never changes, EDF can be classified as a job-level fixed-priority scheduling
    - You’d better to avoid using the term “priority” for EDF since there is no explicit definition

# Earliest-Deadline-First Scheduling


- If an algorithm schedulable  $\leftrightarrow$  feasible
  - It can be referred to as a universal scheduling algorithm
- What is the **universal scheduling algorithm** for periodic and preemptive uniprocessor systems?
  - EDF!

# Earliest-Deadline-First Scheduling


- Example

*depended on*

Not scheduable by RMS



Schedulable by EDF




# Earliest-Deadline-First Scheduling

- ***Theorem:*** With EDF, there is no idle time before an overflow
- **Observation:** A very strong statement that implies optimality of EDF in terms of schedulability



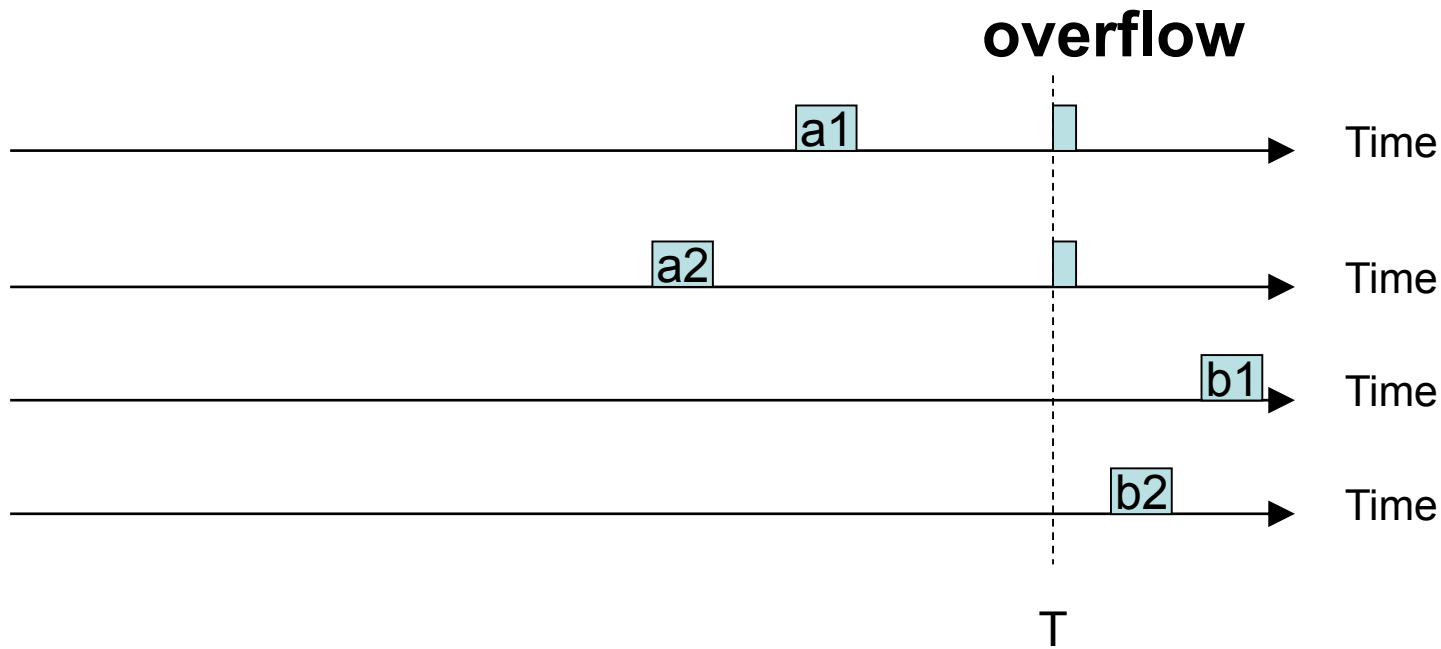
# Proof



- Suppose that there is an overflow at time  $t_3$ , and the processor idles between  $t_1$  and  $t_2$
- If we move "a" forward to be aligned with  $t_2$ , the overflow would occur earlier than it was (i.e., at or before  $t_3$ )
  - That is because EDF's discipline: moving forward means promoting the urgency of  $T_1$ 's jobs
- By repeating the above action, jobs a, b, and c can be aligned at  $t_2$ 
  - → that contradicts the assumption! From  $t_2$  on, there is no idle until the overflow

# Earliest-Deadline-First Scheduling

- ***Theorem:*** A set of tasks is schedulable by EDF if and only if its total CPU utilization is no more than 1
- Observation:  $\rightarrow$  is easy,  $\leftarrow$  requires some reasoning similar to the proof of the last theorem



→: suppose that  $U \leq 1$  but the system is not schedulable by EDF

- Suppose that there is an overflow at time T
  - Jobs a's have deadlines at time T
  - Job b's have deadlines after time T

• Case A: none of job b's is executed before T

- The total computational demand between  $[0, T]$  is

$$C_1(\lfloor T/P_1 \rfloor) + C_2(\lfloor T/P_2 \rfloor) + \dots + C_n(\lfloor T/P_n \rfloor)$$

- Since there is no idle before an overflow


$$C_1(\lfloor T/P_1 \rfloor) + C_2(\lfloor T/P_2 \rfloor) + \dots + C_n(\lfloor T/P_n \rfloor) > T$$

- That implies  $U > 1$

• → ←

**Boxes: “arrival” times!!**

**overflow**



Suppose b2 is  
being executed  
here...

Case B: some of job b's are executed before T

- Because an overflow occurs at T, the violated jobs must be a's
  - Right before T, there must be some job a's being executed
  - Let in  $[T', T]$  there is no job b's being executed
- Just before T', some of b's is being executed! (the definition of T')
  - It means that all jobs have deadlines  $\leq T$  and arrive before T' have completed before T'
- Back to  $[T', T]$ , the total computation demand is no less than
 
$$C_1(\lfloor T - T'/P_1 \rfloor) + C_2(\lfloor T - T'/P_2 \rfloor) + \dots + C_n(\lfloor T - T'/P_n \rfloor)$$
  - Because there is deadline violations, so
 
$$C_1(\lfloor T - T'/P_1 \rfloor) + C_2(\lfloor T - T'/P_2 \rfloor) + \dots + C_n(\lfloor T - T'/P_n \rfloor) \geq T - T'$$

# Earliest-Deadline-First Scheduling

- Summary
  - A universal scheduling algorithm for real-time periodic tasks
  - Urgency of tasks is dynamic
    - But static for jobs
  - Job-level fixed-priority scheduling

# Comparison

	RMS	EDF
Optimality	Optimal for fixed-priority scheduling	Universal
Schedulability test	Exact test is slow (PP), conservative tests are adopted	$O(n)$ for exact test
Algorithm time complexity	$O(1)$ job insertion is possible	Both job insertion and dispatch take $O(\log n)$ time
Overload survivability	High and predictable	Low and unmanageable**
Responsiveness	High priority tasks always have shorter response time	Non-intuitive to reach conclusions
Ease of implementation	Pretty simple	Relatively complicated
Run-time overheads (like preemption)	Low	High

Is it true?